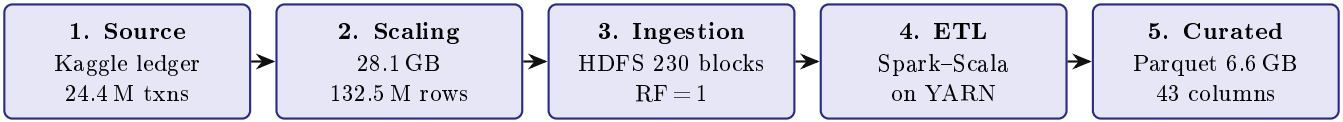
1. Source
Kaggle ledger
24.4 M txns
2. Scaling
28.1 GB
132.5 M rows
3. Ingestion
HDFS 230 blocks
RF = 1
4. ETL
Spark–Scala
on YARN
5. Curated
Parquet 6.6 GB
43 columns